

Introduction

A conditional distribution describes how one random variable behaves given a specific value (or values) of another. Conditional expectations are what regression models estimate; conditional variances quantify residual uncertainty. Together they are the language of predictive statistics.

Prerequisites

Joint distributions, marginal distributions, conditional probability.

Theory

For continuous joint density $f_{X,Y}$ with marginal $f_Y > 0$, the **conditional density of X given $Y = y$** is

$$f_{X|Y}(x | y) = \frac{f_{X,Y}(x, y)}{f_Y(y)}.$$

Analogous expression for discrete variables using PMFs.

Conditional expectation of X given $Y = y$:

$$E[X | Y = y] = \int x f_{X|Y}(x | y) dx.$$

As a function of y , $E[X | Y]$ is itself a random variable.

Tower property (law of total expectation):

$$E[X] = E[E[X | Y]].$$

Conditional variance: $\text{Var}(X | Y = y) = E[X^2 | Y = y] - (E[X | Y = y])^2$.

Law of total variance: $\text{Var}(X) = E[\text{Var}(X | Y)] + \text{Var}(E[X | Y])$. Within-group variance plus between-group variance.

For a bivariate normal with correlation ρ : $X | Y = y \sim \mathcal{N}(\mu_X + \rho\sigma_X\sigma_Y^{-1}(y - \mu_Y), \sigma_X^2(1 - \rho^2))$.

Assumptions

$f_Y(y) > 0$ (or the conditional probability is defined via an alternate limiting procedure); $E[|X|] < \infty$ for the conditional expectation to exist.

R Implementation

```
library(mvtnorm)
set.seed(2026)

# Bivariate normal
mu <- c(0, 0); Sigma <- matrix(c(1, 0.7, 0.7, 1), 2, 2)
rho <- Sigma[1, 2] / sqrt(prod(diag(Sigma)))
XY <- rmvnorm(5e4, mean = mu, sigma = Sigma)

# Conditional on Y near y0, distribution of X
y0 <- 1.0; window <- 0.05
X_given_Y <- XY[abs(XY[, 2] - y0) < window, 1]
```

```

c(n_window      = length(X_given_Y),
  emp_mean      = mean(X_given_Y),
  theor_mean    = mu[1] + rho * (y0 - mu[2]),
  emp_sd        = sd(X_given_Y),
  theor_sd      = sqrt(Sigma[1, 1] * (1 - rho^2)))

# Tower property:  $E[X] = E[E[X | Y]]$ 
c(E_X = mean(XY[, 1]),
  tower = mean(mu[1] + rho * (XY[, 2] - mu[2])))

# Law of total variance
var_total <- var(XY[, 1])
var_within <- Sigma[1, 1] * (1 - rho^2)
var_between <- rho^2 * Sigma[1, 1]
c(total = var_total,
  sum_of_parts = var_within + var_between)

```

Output & Results

n_window	emp_mean	theor_mean	emp_sd	theor_sd
4094	0.693	0.700	0.708	0.714

E_X	tower
0.003	-0.001

total	sum_of_parts
1.000	1.000

Empirical conditional moments match theoretical; tower property and law of total variance hold exactly (up to Monte Carlo error).

Interpretation

Regression is the enterprise of estimating $E[Y | X = x]$ as a function of x . When a manuscript reports “for every 1-unit increase in x , y increases by β ”, it is interpreting a slope of the conditional expectation.

Practical Tips

- Conditional distributions require the conditioning event to have positive probability (or a limiting argument for continuous variables).
- $E[Y | X]$ is a function of X and therefore a random variable; $E[Y | X = x]$ is a specific number.
- The conditional expectation is the best predictor of Y from X under squared-error loss (Hilbert projection).
- Jensen’s inequality extends to conditional expectations: $E[g(X) | Y] \geq g(E[X | Y])$ for convex g .
- In high-dimensional conditioning, direct computation is intractable; regression models (linear, generalised, spline-based) estimate $E[Y | X]$ directly.

For Reviewers

What to look for in a paper using this method.

- Common misapplications.
- Diagnostics that should be reported but often aren’t.
- Red flags in tables and figures.
- What to verify.

- **What an adequate Methods paragraph must contain.**

See also — labs in this chapter

- Probability, conditional probability, Bayes' theorem
- Discrete distributions: Bernoulli, binomial, Poisson
- Continuous distributions and Q-Q plots