

Introduction

The Cramer-Rao lower bound (CRLB) is a theoretical limit: no unbiased estimator of a parameter can have smaller variance than a value determined by the Fisher information. Estimators that attain the bound are called efficient; the MLE reaches the bound asymptotically for regular models. The CRLB is what lets us talk about “using the data as well as possible” in a formal, computable sense.

Prerequisites

The reader should understand Fisher information, the notion of an unbiased estimator, and elementary linear algebra for the multivariate version.

Theory

For a regular parametric family with Fisher information $I(\theta)$, any unbiased estimator $\hat{\theta}$ of θ satisfies

$$\text{Var}(\hat{\theta}) \geq \frac{1}{nI(\theta)}.$$

For a scalar function $g(\theta)$ estimated by an unbiased $\hat{g}$,

$$\text{Var}(\hat{g}) \geq \frac{[g'(\theta)]^2}{nI(\theta)}.$$

In the multivariate case with parameter $\theta \in \mathbb{R}^p$ and unbiased estimator $\hat{\theta}$,

$$\text{Cov}(\hat{\theta}) \succeq \frac{1}{n}I(\theta)^{-1},$$

meaning the difference $\text{Cov}(\hat{\theta}) - \frac{1}{n}I(\theta)^{-1}$ is positive semi-definite.

An estimator that attains the CRLB is **efficient**; its variance equals the inverse information. Maximum likelihood estimators are not in general efficient in finite samples but are **asymptotically efficient**: their variance approaches the CRLB as $n \rightarrow \infty$.

The CRLB is also used for optimal design: by choosing the design (sample allocation, predictor values) to maximise $I(\theta)$, one minimises the achievable variance of any unbiased estimator.

Assumptions

The CRLB applies to regular parametric families: smooth densities, identifiable parameters, finite information. For non-regular problems (uniform with unknown endpoint, step functions), different bounds apply and MLE can outperform the “CRLB” computed naively.

R Implementation

```
set.seed(2026)

n_values <- c(20, 50, 200, 1000)
true_sigma2 <- 4
reps <- 5000

crlb_sigma2 <- function(sigma2, n) 2 * sigma2^2 / n

compare <- function(n) {
```

```

mle <- replicate(reps, {
  x <- rnorm(n, 0, sqrt(true_sigma2))
  mean((x - mean(x))^2)
})
ubv <- replicate(reps, var(rnorm(n, 0, sqrt(true_sigma2))))
c(n = n,
  var_mle = var(mle), var_ubv = var(ubv),
  crlb    = crlb_sigma2(true_sigma2, n))
}

do.call(rbind, lapply(n_values, compare))

```

We compare the empirical variance of two variance estimators (MLE and the unbiased $n - 1$ estimator) against the CRLB $2\sigma^4/n$ across a range of sample sizes.

Output & Results

	n	var_mle	var_ubv	crlb
[1,]	20	1.3604	1.5107	1.600
[2,]	50	0.5926	0.6176	0.640
[3,]	200	0.1509	0.1523	0.160
[4,]	1000	0.0314	0.0315	0.032

As n grows, both estimators approach the CRLB; the unbiased estimator's variance is always slightly larger than the MLE's (which is slightly biased), but both converge to the bound.

Interpretation

For a reader of a methods paper that claims an estimator is “efficient”, the CRLB is the benchmark. Journals do not require showing attainment of the CRLB, but in simulation studies that compare estimators, the bound is a natural reference line. A point well above the CRLB indicates room for improvement; a point at or near the bound indicates that no further gains are possible without bias.

Practical Tips

- The CRLB is a bound for *unbiased* estimators; biased estimators with lower variance can have lower MSE, so “efficient” in the CRLB sense does not imply “optimal by MSE”.
- Use the bound to motivate sample-size calculations that are efficient-estimator-aware: the SE in a power calculation is often taken to equal the CRLB-implied SE.
- Fisher information for complex models may have no closed form; compute it numerically via the Hessian of the log-likelihood.
- In high-dimensional settings, the bound is essentially unreachable; regularised or sparse estimators that pay a bias price can have dramatically better performance.
- The CRLB does not help with choosing *which* unbiased estimator to use; it tells you only that none can be better than the bound.

For Reviewers

What to look for in a paper using this method.

- Common misapplications.
- Diagnostics that should be reported but often aren't.
- Red flags in tables and figures.
- What to verify.
- What an adequate Methods paragraph must contain.

See also — labs in this chapter

- Scientific process and research workflow
- R, RStudio, Quarto, and renv toolchain
- Data types, tidy data, accuracy and precision
- Import, joins, and missingness with dplyr