

# The Gap Statistic

2026-04-17

## Introduction

The gap statistic, introduced by Tibshirani, Walther, and Hastie in 2001, is a principled way to choose the number of clusters  $k$  in any partitioning method. Where the elbow method relies on visual judgement and the silhouette method on per-observation neighbour comparisons, the gap statistic compares the within-cluster dispersion of the data to what would be expected under a null reference distribution with no cluster structure. The point at which the data deviate most from the null is the optimal  $k$ , with bootstrap-based standard errors that quantify the certainty of the choice.

## Prerequisites

A working understanding of clustering algorithms, within-cluster sum of squares, and bootstrap resampling for standard error estimation.

## Theory

For each candidate  $k$ , define  $W_k$  as the log of the within-cluster dispersion (sum of squared distances to centroids) on the actual data after running  $k$ -means or another partitioning algorithm. Generate  $B$  datasets from a null reference — uniformly distributed over the data’s bounding box, or over its principal-components box — and compute the average within-cluster dispersion  $\mathbb{E}[W_k^*]$  under the null. The gap statistic is

$$\text{Gap}(k) = \mathbb{E}[W_k^*] - W_k.$$

The optimal  $k$  is the smallest value satisfying  $\text{Gap}(k) \geq \text{Gap}(k+1) - s_{k+1}$ , where  $s_{k+1}$  is the standard deviation of the null reference at  $k+1$  (the so-called “first-SE-max” criterion). This rule rewards parsimonious choices: it picks the smallest  $k$  that is not significantly worse than larger values.

## Assumptions

A meaningful distance metric on the data; the null reference distribution is sensible (uniform over the bounding box is reasonable for low-to-moderate dimensions; PC-aligned bounding boxes are better in high dimensions).

## R Implementation

```
library(cluster); library(factoextra)

X <- scale(iris[, 1:4])

set.seed(2026)
gap <- clusGap(X, FUN = kmeans, nstart = 25, K.max = 10, B = 100)
print(gap, method = "firstSEmax")

fviz_gap_stat(gap)
```

## Output & Results

`clusGap()` returns a matrix with  $W_k$ ,  $\mathbb{E}[W_k^*]$ , the gap value, and its bootstrap standard error for each  $k$ . `print(..., method = "firstSEmax")` applies Tibshirani's selection rule to identify the optimal  $k$ . `fviz_gap_stat()` plots the gap curve with error bars.

## Interpretation

A reporting sentence: "The gap statistic, computed with  $B = 100$  bootstrap replicates against a uniform reference, identified  $k = 3$  as the optimal number of clusters using the firstSEmax criterion ( $\text{Gap}(3) = 0.59$ ,  $\text{Gap}(4) - s_4 = 0.55$ ); this matches the underlying species structure in the iris data." Report the criterion explicitly (firstSEmax, Tibs2001SEmax, globalmax); they sometimes disagree.

## Practical Tips

- Use  $B = 100$  to 500 bootstrap replicates; fewer can produce unstable optima, more is rarely worth the runtime cost.
- The default uniform-bounding-box reference works well in low dimensions; in high dimensions, switch to a PC-aligned reference (`spaceH0 = "scaledPCA"`).
- Combine the gap statistic with elbow and silhouette methods; concordance across all three is the strongest argument for a chosen  $k$ .
- Gap performs best on well-separated clusters; for overlapping or hierarchical structures it can favour very large  $k$  or fail to find a clear optimum.
- For very large datasets, subsample before applying the gap statistic; the bootstrap step otherwise dominates runtime.
- Selection criteria differ: `firstSEmax` (Tibshirani's parsimonious choice), `globalmax` (the absolute maximum, often larger), and `Tibs2001SEmax` (the original 2001 rule). Quote the criterion in any report.

## R Packages Used

`cluster` for `clusGap()` and the underlying clustering algorithms; `factoextra` for `fviz_gap_stat()` and `fviz_nbclust()` ggplot-style visualisation; `NbClust` for 30+ alternative cluster-validation indices in one call.

## For Reviewers

What to look for in a paper using this method.

- **Common misapplications.**
- **Diagnostics that should be reported but often aren't.**
- **Red flags in tables and figures.**
- **What to verify.**
- **What an adequate Methods paragraph must contain.**

## See also — labs in this chapter

- PCA, factor analysis, CCA, LDA
- Clustering: k-means, hierarchical, model-based
- UMAP and t-SNE