

# Kolmogorov-Smirnov Test

2026-04-17

## Introduction

The Kolmogorov-Smirnov (KS) test compares an empirical distribution function (ECDF) against either a hypothesised theoretical CDF (one-sample test) or another empirical distribution (two-sample test). Its test statistic is the supremum of the absolute difference between the two CDFs, a measure of the largest vertical gap on a CDF plot. KS is the prototypical CDF-based goodness-of-fit test, valued for being fully non-parametric and distribution-free in the one-sample case (the null distribution of  $D$  does not depend on the specific  $F_0$  for continuous distributions). It is widely used as a quick distributional diagnostic and as a flexible two-sample test sensitive to differences in location, scale, and shape.

## Prerequisites

A working understanding of empirical CDFs, the Glivenko-Cantelli theorem, and the difference between testing a fully specified hypothesised distribution and testing a parametric family with parameters estimated from the data.

## Theory

For the **one-sample** test with sample  $X_1, \dots, X_n$  and hypothesised CDF  $F_0$ ,

$$D = \sup_x |\hat{F}_n(x) - F_0(x)|;$$

under  $H_0 : X \sim F_0$ ,  $\sqrt{n} \cdot D$  has the Kolmogorov limiting distribution. For the **two-sample** test with empirical CDFs  $\hat{F}_n$  and  $\hat{G}_m$ ,

$$D_{n,m} = \sup_x |\hat{F}_n(x) - \hat{G}_m(x)|.$$

KS is distribution-free in the continuous one-sample case: the null distribution of  $D$  depends only on the sample size, not the specific  $F_0$ . When parameters of  $F_0$  are estimated from the same data, the nominal  $p$ -values are anti-conservative; the Lilliefors correction provides the appropriate adjustment for testing Normality with estimated mean and SD.

## Assumptions

The underlying distribution is continuous (ties produce warnings and approximate  $p$ -values), observations are independent and identically distributed, and parameters of  $F_0$  are either fully specified or appropriately corrected for if estimated from the sample.

## R Implementation

```
set.seed(2026)

x <- rlnorm(100, meanlog = 0, sdlog = 0.8)
```

```

ks.test(x, "pnorm", mean = mean(x), sd = sd(x))

z <- scale(x)
ks.test(z, "pnorm")

x1 <- rnorm(50, 0, 1)
x2 <- rnorm(60, 0.4, 1.2)
ks.test(x1, x2)

library(nortest)
lillie.test(x)

```

## Output & Results

The script demonstrates the four standard use cases. The one-sample KS against a Normal with sample-estimated parameters rejects the null but the  $p$ -value is anti-conservative; the Lilliefors-corrected variant gives the proper rejection. The two-sample KS detects the location and scale shift between  $X_1$  and  $X_2$  ( $D = 0.31$ ,  $p = 0.007$ ), illustrating the test's sensitivity to general distributional differences.

## Interpretation

A reporting sentence: “The Lilliefors-corrected Kolmogorov-Smirnov test rejected Normality of the residuals from the regression model ( $D = 0.16$ ,  $p < 0.001$ ), confirming the visual impression of right-skew in the Q-Q plot. A log transformation was applied before re-fitting; the residuals on the transformed scale passed the same test ( $D = 0.05$ ,  $p = 0.21$ ). The two-sample KS comparing transformed-residual distributions across treatment arms detected no significant difference in shape ( $D = 0.08$ ,  $p = 0.52$ ).” Always pair KS with a Q-Q plot.

## Practical Tips

- KS is generally less powerful than Shapiro-Wilk for testing Normality in small to moderate samples; prefer Shapiro-Wilk as the default Normality test and use KS when the hypothesised distribution is non-Normal or when a two-sample distributional comparison is needed.
- Ties in the sample invalidate the strict KS  $p$ -value and R issues a warning; prefer Anderson-Darling or the Cramér-von Mises test with discretised or rounded data, which handle ties more gracefully.
- The Lilliefors correction must be used when testing Normality with sample-estimated mean and SD; the standard `ks.test()` against `pnorm` with sample estimates gives anti-conservative  $p$ -values that reject too rarely.
- The two-sample KS detects any distributional difference — location, scale, shape — and is therefore a useful omnibus diagnostic when the analyst is unsure which aspect of the distribution might differ; it is generally less powerful than the Mann-Whitney  $U$  test for pure location shifts.
- For very large samples, KS will reject almost any parametric fit because real data have rounding, ceiling effects, and other minor non-Normalities; use effect-size summaries and Q-Q plots in large samples rather than relying on the  $p$ -value alone.
- The Anderson-Darling test gives more weight to the tails of the distribution than KS does; for tail-sensitive applications (extreme-value detection, financial risk), Anderson-Darling is often preferred.

## R Packages Used

Base R `ks.test()` for one- and two-sample Kolmogorov-Smirnov; `nortest::lillie.test()` for the Lilliefors-corrected Normality variant; `nortest::ad.test()` and `nortest::cvm.test()` for Anderson-Darling and Cramér-von Mises alternatives; `dgof::ks.test()` for discrete-data extensions; `car::qqPlot()` for graphical complements to the formal test.

## For Reviewers

What to look for in a paper using this method.

- Common misapplications.
- Diagnostics that should be reported but often aren't.
- Red flags in tables and figures.
- What to verify.
- What an adequate Methods paragraph must contain.

See also — labs in this chapter

- Populations, samples, and the central limit theorem
- Bootstrap and permutation tests
- Maximum likelihood estimation
- One-sample t-test and one-proportion test
- Hypothesis testing, p-values, type I/II errors
- Two-sample and paired t-tests
- Two proportions, chi-square, risk and odds
- Pearson and Spearman correlation
- Non-parametric tests