

Imaging and sequence models intro

Workflow labs use the variant template: Goal → Approach → Execution → Check → Report.

Learning objectives

- Describe how a 2-D convolution extracts spatial features from an image.
- Describe how a 1-D convolution or a recurrent layer processes a sequence.
- Recognise when to use a CNN, RNN, or transformer for biomedical data.

Prerequisites

A neural-network training loop in `torch`.

Background

Convolutional neural networks (CNNs) exploit the spatial locality of image data: a small filter is slid across the image, producing a feature map for each filter. Stacking convolutional layers lets the network learn increasingly abstract features, from edges to textures to whole tissue structures. In biomedical imaging — histopathology, radiology, dermatology — CNNs underpin the majority of deep-learning applications.

For sequence data — gene sequences, electronic health record time series, wearable signals — the relevant operations are 1-D convolutions and attention. Transformers have become the dominant architecture for long sequences, but 1-D CNNs and gated recurrent networks remain competitive on shorter inputs and are easier to train. This lab illustrates the *shape* of the computation with a tiny hand-built 2-D convolution and a tiny 1-D convolution; the real training happens with frameworks that we mark `eval: false`.

The conceptual point: image and sequence models share the training tabular `torch` training loop. What differs is the architecture, the data pipeline, and the inductive biases (locality for CNNs, ordered context for RNNs and attention).

Setup

```
library(tidyverse)
set.seed(42)
theme_set(theme_minimal(base_size = 12))
```

1. Goal

Show manually computed convolutions on a small image and sequence, then sketch the `torch` code that would perform the same operations inside a trained model.

2. Approach

Create a toy 16×16 image with a bright square.

```
img[5:10, 5:10] <- 1
img <- img + matrix(rnorm(256, 0, 0.1), 16, 16)

image(t(img[16:1, ]), col = grey.colors(100), main = "toy image")
```

3. Execution

A manual 3×3 edge filter convolved with the image.

```
conv2d <- function(img, k) {  
  kh <- nrow(k); kw <- ncol(k)  
  out <- matrix(0, nrow(img) - kh + 1, ncol(img) - kw + 1)  
  for (i in seq_len(nrow(out))) for (j in seq_len(ncol(out)))  
    out[i, j] <- sum(img[i:(i + kh - 1), j:(j + kw - 1)] * k)  
  out  
}  
feat <- conv2d(img, kernel)  
image(t(feat[nrow(feat):1, ]), col = grey.colors(100),  
      main = "edge-filter feature map")
```

A 1-D convolution on a toy sequence.

```
conv1d <- function(x, k) {  
  out <- numeric(length(x) - length(k) + 1)  
  for (i in seq_along(out)) out[i] <- sum(x[i:(i + length(k) - 1)] * k)  
  out  
}  
feat1d <- conv1d(seq1, c(-1, 1))  
tibble(i = seq_along(feat1d), v = feat1d) |>  
  ggplot(aes(i, v)) + geom_line() +  
  labs(x = "position", y = "edge response")
```

A tiny CNN in torch (sketch only).

```
library(torch)  
cnn <- nn_sequential(  
  nn_conv2d(1, 8, kernel_size = 3, padding = 1),  
  nn_relu(),  
  nn_max_pool2d(2),  
  nn_conv2d(8, 16, kernel_size = 3, padding = 1),  
  nn_relu(),  
  nn_max_pool2d(2),  
  nn_flatten(),  
  nn_linear(16 * 4 * 4, 10)  
)
```

A 1-D convolutional sequence model sketch.

```
library(torch)  
seq_model <- nn_sequential(  
  nn_conv1d(1, 8, kernel_size = 5, padding = 2),  
  nn_relu(),  
  nn_adaptive_avg_pool1d(1),  
  nn_flatten(),  
  nn_linear(8, 1)  
)
```

4. Check

Inspect the convolution output shape and magnitude.

```
summary(as.numeric(feat))
```

5. Report

A 3×3 edge filter applied to a 16×16 toy image produced a feature map of shape `paste(dim(feat), collapse = "x")`, with responses concentrated at the boundary of the bright square. Analogous 1-D convolutions applied to a binary step sequence produce sharp derivative-like responses at the transitions.

Real CNNs and sequence models compose many such filters, learned from data. The lab establishes the building block so the architectures in the literature read as stacks of familiar operations, not black boxes.

Common pitfalls

- Applying a CNN to tabular data; locality inductive bias is wasted.
- Forgetting input channel and batch dimensions in `torch`.
- Comparing accuracy of sequence models without specifying the tokenisation used.

Further reading

- LeCun Y, Bengio Y, Hinton G (2015), *Deep learning*.
- Vaswani A et al. (2017), *Attention is all you need*.

Session info

See also — chapter index

- Methods in this chapter
- Find your method (Chapter 0)